

Jacobian Descent For Multi-Objective Optimization

Valérian Rey — Simplex Lab
valerian.rey@gmail.com
Joint work with Pierre Quinton

February 11, 2026



github.com/TorchJD/torchjd



torchjd.org



arxiv.org/pdf/2406.16232

Table of contents

Theory

- Partial order

- Pareto optimality

- Jacobian descent

- Unconflicting projection of gradients (UPGrad)

Applications and experimentation

- Optimization trajectories

- Per-sample JD

- Multi-task learning

Conclusion

Theory

Partial order

Define the partial order $(\mathbb{R}^m, \preceq)$, for $\mathbf{u}, \mathbf{v} \in \mathbb{R}^m$, as

$$\begin{aligned} \mathbf{u} \preceq \mathbf{v} \\ \Leftrightarrow \quad u_i \leq v_i \quad \forall i \in [m] \end{aligned}$$

Partial order

Define the partial order $(\mathbb{R}^m, \preceq)$, for $\mathbf{u}, \mathbf{v} \in \mathbb{R}^m$, as

$$\begin{aligned} \mathbf{u} \preceq \mathbf{v} \\ \Leftrightarrow u_i \leq v_i \quad \forall i \in [m] \end{aligned}$$

Examples:

$$\blacktriangleright \begin{bmatrix} 2 \\ 1 \end{bmatrix} \preceq \begin{bmatrix} 2 \\ 2 \end{bmatrix}$$

Partial order

Define the partial order $(\mathbb{R}^m, \preceq)$, for $\mathbf{u}, \mathbf{v} \in \mathbb{R}^m$, as


$$\begin{aligned} \mathbf{u} \preceq \mathbf{v} \\ \Leftrightarrow u_i \leq v_i \quad \forall i \in [m] \end{aligned}$$

Examples:

▶ $\begin{bmatrix} 2 \\ 1 \end{bmatrix} \preceq \begin{bmatrix} 2 \\ 2 \end{bmatrix}$


▶ $\begin{bmatrix} 2 \\ 1 \end{bmatrix}$ and $\begin{bmatrix} 1 \\ 2 \end{bmatrix}$ are not comparable with this partial order.

Pareto optimality




Pareto optimality

Pareto set: X^*




Pareto optimality

Pareto set: X^*
Pareto Front: $f(X^*)$



Pareto optimality

Pareto set: X^*
Pareto Front: $f(X^*)$



$$\mathbf{x}^* \in X^* \quad \Leftrightarrow \quad \forall \mathbf{y} \in \mathbb{R}^n, \mathbf{f}(\mathbf{y}) \preceq \mathbf{f}(\mathbf{x}^*) \rightarrow \mathbf{f}(\mathbf{y}) = \mathbf{f}(\mathbf{x}^*)$$

Jacobian descent

Goal: minimize $\mathbf{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ according to the $\preceq$ partial order.

Jacobian descent

Goal: minimize $\mathbf{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ according to the $\preceq$ partial order. For any $\mathbf{x} \in \mathbb{R}^n$, the Jacobian of $\mathbf{f}$ at $\mathbf{x}$, is

$$\mathcal{J}\mathbf{f}(\mathbf{x}) = \begin{bmatrix} \nabla f_1(\mathbf{x})^\top \\ \nabla f_2(\mathbf{x})^\top \\ \vdots \\ \nabla f_m(\mathbf{x})^\top \end{bmatrix} = \begin{bmatrix} \frac{\partial}{\partial x_1} f_1(\mathbf{x}) & \frac{\partial}{\partial x_2} f_1(\mathbf{x}) & \cdots & \frac{\partial}{\partial x_n} f_1(\mathbf{x}) \\ \frac{\partial}{\partial x_1} f_2(\mathbf{x}) & \frac{\partial}{\partial x_2} f_2(\mathbf{x}) & \cdots & \frac{\partial}{\partial x_n} f_2(\mathbf{x}) \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial}{\partial x_1} f_m(\mathbf{x}) & \frac{\partial}{\partial x_2} f_m(\mathbf{x}) & \cdots & \frac{\partial}{\partial x_n} f_m(\mathbf{x}) \end{bmatrix}$$

Jacobian descent

Goal: minimize $\mathbf{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ according to the $\preceq$ partial order. For any $\mathbf{x} \in \mathbb{R}^n$, the Jacobian of $\mathbf{f}$ at $\mathbf{x}$, is

$$\mathcal{J}\mathbf{f}(\mathbf{x}) = \begin{bmatrix} \nabla f_1(\mathbf{x})^\top \\ \nabla f_2(\mathbf{x})^\top \\ \vdots \\ \nabla f_m(\mathbf{x})^\top \end{bmatrix} = \begin{bmatrix} \frac{\partial}{\partial x_1} f_1(\mathbf{x}) & \frac{\partial}{\partial x_2} f_1(\mathbf{x}) & \cdots & \frac{\partial}{\partial x_n} f_1(\mathbf{x}) \\ \frac{\partial}{\partial x_1} f_2(\mathbf{x}) & \frac{\partial}{\partial x_2} f_2(\mathbf{x}) & \cdots & \frac{\partial}{\partial x_n} f_2(\mathbf{x}) \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial}{\partial x_1} f_m(\mathbf{x}) & \frac{\partial}{\partial x_2} f_m(\mathbf{x}) & \cdots & \frac{\partial}{\partial x_n} f_m(\mathbf{x}) \end{bmatrix}$$

For a small enough update $\mathbf{y} \in \mathbb{R}^n$, Taylor's theorem states

$$\mathbf{f}(\mathbf{x} + \mathbf{y}) \approx \mathbf{f}(\mathbf{x}) + \mathcal{J}\mathbf{f}(\mathbf{x}) \cdot \mathbf{y}$$

Jacobian descent

Goal: minimize $\mathbf{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ according to the $\preceq$ partial order. For any $\mathbf{x} \in \mathbb{R}^n$, the Jacobian of $\mathbf{f}$ at $\mathbf{x}$, is

$$\mathcal{J}\mathbf{f}(\mathbf{x}) = \begin{bmatrix} \nabla f_1(\mathbf{x})^\top \\ \nabla f_2(\mathbf{x})^\top \\ \vdots \\ \nabla f_m(\mathbf{x})^\top \end{bmatrix} = \begin{bmatrix} \frac{\partial}{\partial x_1} f_1(\mathbf{x}) & \frac{\partial}{\partial x_2} f_1(\mathbf{x}) & \cdots & \frac{\partial}{\partial x_n} f_1(\mathbf{x}) \\ \frac{\partial}{\partial x_1} f_2(\mathbf{x}) & \frac{\partial}{\partial x_2} f_2(\mathbf{x}) & \cdots & \frac{\partial}{\partial x_n} f_2(\mathbf{x}) \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial}{\partial x_1} f_m(\mathbf{x}) & \frac{\partial}{\partial x_2} f_m(\mathbf{x}) & \cdots & \frac{\partial}{\partial x_n} f_m(\mathbf{x}) \end{bmatrix}$$

For a small enough update $\mathbf{y} \in \mathbb{R}^n$, Taylor's theorem states

$$\mathbf{f}(\mathbf{x} + \mathbf{y}) \approx \mathbf{f}(\mathbf{x}) + \mathcal{J}\mathbf{f}(\mathbf{x}) \cdot \mathbf{y}$$

We want $\mathbf{f}(\mathbf{x} + \mathbf{y}) \preceq \mathbf{f}(\mathbf{x})$.

Jacobian descent

Goal: minimize $\mathbf{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ according to the $\preceq$ partial order. For any $\mathbf{x} \in \mathbb{R}^n$, the Jacobian of $\mathbf{f}$ at $\mathbf{x}$, is

$$\mathcal{J}\mathbf{f}(\mathbf{x}) = \begin{bmatrix} \nabla f_1(\mathbf{x})^\top \\ \nabla f_2(\mathbf{x})^\top \\ \vdots \\ \nabla f_m(\mathbf{x})^\top \end{bmatrix} = \begin{bmatrix} \frac{\partial}{\partial x_1} f_1(\mathbf{x}) & \frac{\partial}{\partial x_2} f_1(\mathbf{x}) & \cdots & \frac{\partial}{\partial x_n} f_1(\mathbf{x}) \\ \frac{\partial}{\partial x_1} f_2(\mathbf{x}) & \frac{\partial}{\partial x_2} f_2(\mathbf{x}) & \cdots & \frac{\partial}{\partial x_n} f_2(\mathbf{x}) \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial}{\partial x_1} f_m(\mathbf{x}) & \frac{\partial}{\partial x_2} f_m(\mathbf{x}) & \cdots & \frac{\partial}{\partial x_n} f_m(\mathbf{x}) \end{bmatrix}$$

For a small enough update $\mathbf{y} \in \mathbb{R}^n$, Taylor's theorem states

$$\mathbf{f}(\mathbf{x} + \mathbf{y}) \approx \mathbf{f}(\mathbf{x}) + \mathcal{J}\mathbf{f}(\mathbf{x}) \cdot \mathbf{y}$$

We want $\mathbf{f}(\mathbf{x} + \mathbf{y}) \preceq \mathbf{f}(\mathbf{x})$. We select $\mathbf{y} = -\eta \mathcal{A}(\mathcal{J}\mathbf{f}(\mathbf{x})) \in \mathbb{R}^n$.

Jacobian descent

The mapping $\mathcal{A} : \mathbb{R}^{m \times n} \rightarrow \mathbb{R}^n$ is called an **aggregator**, it parameterizes **Jacobian descent**:

Jacobian descent

The mapping $\mathcal{A} : \mathbb{R}^{m \times n} \rightarrow \mathbb{R}^n$ is called an **aggregator**, it parameterizes **Jacobian descent**:

Algorithm 1: Jacobian descent with aggregator $\mathcal{A}$


Input: $\mathbf{x} \in \mathbb{R}^n$, $0 < \eta$, $T \in \mathbb{N}$, $\mathcal{A} : \mathbb{R}^{m \times n} \rightarrow \mathbb{R}^n$

for $t \leftarrow 1$ **to** T **do**


$\mathbf{x} \leftarrow \mathbf{x} - \eta \mathcal{A}(\mathcal{J}\mathbf{f}(\mathbf{x}))$

Output: $\mathbf{x}$


Unconflicting projection of gradients (UPGrad)




Unconflicting projection of gradients (UPGrad)




Unconflicting projection of gradients (UPGrad)



Unconflicting projection of gradients (UPGrad)



Unconflicting projection of gradients (UPGrad)




Formal definition of UPGrad

Given $J \in \mathbb{R}^{m \times n}$, let C_J be the cone $\{\mathbf{y} \in \mathbb{R}^n : \mathbf{0} \preceq J\mathbf{y}\}$. $\mathcal{A}$ is non-conflicting if for all $J \in \mathbb{R}^{m \times n}$, $\mathcal{A}(J) \in C_J$.

For any $\mathbf{x} \in \mathbb{R}^n$, the projection of $\mathbf{x}$ onto C_J is

$$\pi_J(\mathbf{x}) = \arg \min_{\mathbf{y} \in C_J} \|\mathbf{x} - \mathbf{y}\|^2$$

$$\mathcal{A}_{\text{UPGrad}}(J) = \frac{1}{m} \sum_{i \in [m]} \pi_J(J^\top \mathbf{e}_i)$$





Applications and experimentation

Optimization trajectories

Goal: Minimize some convex quadratic function $f : \mathbb{R}^2 \rightarrow \mathbb{R}$, with conflicting gradients.

Optimization trajectories

Goal: Minimize some convex quadratic function $f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$, with conflicting gradients.



Optimization trajectories: code

```
x = tensor([0., 0.]) # initial value of x

for i in range(n_iters):
    y = f(x) # shape: [2]
    autojac.backward(y)
    # x now has a .jac field
    autojac.jac_to_grad([x], UPGrad())
    # x now has a .grad field
    optimizer.step()
    optimizer.zero_grad()
```

Full project available at github.com/TorchJD/trajectories

Per-sample JD

Each element of the batch has its own loss.

Per-sample JD

Each element of the batch has its own loss.

```
batch_size = 16
model = Linear(10, 1)

for x, y in dataloader:
    z = model(x).squeeze(dim=1)  # shape: [16]
    losses = loss_fn(z, y)  # shape: [16]
    autograd.backward(losses)
    # parameters now have a .jac field
    autograd.jac_to_grad(model.parameters(), UPGrad())
    # parameters now have a .grad field
    optimizer.step()
    optimizer.zero_grad()
```

Full code example available at torchjd.org/latest/examples/iwrm.


Per-sample JD: results

Batch size: 32


CIFAR-10




Training Loss




Update Similarity




Multi-task learning



Multi-task learning



Multi-task learning



Multi-task learning: code

```
h = Linear(10, 5)
g1 = Linear(5, 1)
g2 = Linear(5, 1)

for x, y1, y2 in dataloader:
    r = h(x)
    z1 = g1(r)
    z2 = g2(r)
    l1 = loss_fn(z1, y1)
    l2 = loss_fn(z2, y2)
    autograd.backward([l1, l2], features=r)
    # backbone parameters now have a .jac field
    # and head parameters have a .grad field
    autograd.backward(backbone.parameters(), UPGrad())
    # All parameters now have a .grad field
    optimizer.step()
    optimizer.zero_grad()
```

Full code example available at [torchjd.org/latest/examples/mtl](https://pytorch.org/latest/examples/mtl).

Conclusion

Many problems are multi-objective.

Conclusion

Many problems are multi-objective.

Dual cone is rarely empty in deep learning.

Conclusion

Many problems are multi-objective.

Dual cone is rarely empty in deep learning.

TorchJD to compute and aggregate Jacobians.

Conclusion

Many problems are multi-objective.

Dual cone is rarely empty in deep learning.


TorchJD to compute and aggregate Jacobians.

Try to compute the inner products between your gradients.

Appendix


Element-wise quadratic function trajectories

Minimize $\mathbf{f} : \begin{bmatrix} x \\ y \end{bmatrix} \mapsto \begin{bmatrix} x^2 \\ y^2 \end{bmatrix}$. It's Jacobian at $\begin{bmatrix} x \\ y \end{bmatrix}$ is $\begin{bmatrix} 2x & 0 \\ 0 & 2y \end{bmatrix}$.



Convex quadratic form trajectories

Some convex quadratic function $\mathbf{f} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$, with conflicting gradients.



Gramian-based JD

$\mathcal{A}_{\text{UPGrad}}(\mathcal{J}\mathbf{f}(\mathbf{x}))$ is a linear combination of the rows of $\mathcal{J}\mathbf{f}(\mathbf{x})$.

Gramian-based JD

$\mathcal{A}_{\text{UPGrad}}(\mathcal{J}\mathbf{f}(\mathbf{x}))$ is a linear combination of the rows of $\mathcal{J}\mathbf{f}(\mathbf{x})$.

The weights only depend on the (small) Gramian
 $\mathcal{G}\mathbf{f}(\mathbf{x}) = \mathcal{J}\mathbf{f}(\mathbf{x}) \cdot \mathcal{J}\mathbf{f}(\mathbf{x})^\top$.

Gramian-based JD

$\mathcal{A}_{\text{UPGrad}}(\mathcal{J}\mathbf{f}(\mathbf{x}))$ is a linear combination of the rows of $\mathcal{J}\mathbf{f}(\mathbf{x})$.

The weights only depend on the (small) Gramian
 $\mathcal{G}\mathbf{f}(\mathbf{x}) = \mathcal{J}\mathbf{f}(\mathbf{x}) \cdot \mathcal{J}\mathbf{f}(\mathbf{x})^\top$.

We can compute $\mathcal{G}\mathbf{f}(\mathbf{x})$ directly, extract the weights, combine the losses, and do a step of gradient descent.

Gramian-based JD

$\mathcal{A}_{\text{UPGrad}}(\mathcal{J}\mathbf{f}(\mathbf{x}))$ is a linear combination of the rows of $\mathcal{J}\mathbf{f}(\mathbf{x})$.

The weights only depend on the (small) Gramian
 $\mathcal{G}\mathbf{f}(\mathbf{x}) = \mathcal{J}\mathbf{f}(\mathbf{x}) \cdot \mathcal{J}\mathbf{f}(\mathbf{x})^\top$.

We can compute $\mathcal{G}\mathbf{f}(\mathbf{x})$ directly, extract the weights, combine the losses, and do a step of gradient descent.

We have developed the `autogram` engine to compute $\mathcal{G}\mathbf{f}(\mathbf{x})$ efficiently (work in progress).

Gramian-based JD

$\mathcal{A}_{\text{UPGrad}}(\mathcal{J}\mathbf{f}(\mathbf{x}))$ is a linear combination of the rows of $\mathcal{J}\mathbf{f}(\mathbf{x})$.

The weights only depend on the (small) Gramian
 $\mathcal{G}\mathbf{f}(\mathbf{x}) = \mathcal{J}\mathbf{f}(\mathbf{x}) \cdot \mathcal{J}\mathbf{f}(\mathbf{x})^\top$.

We can compute $\mathcal{G}\mathbf{f}(\mathbf{x})$ directly, extract the weights, combine the losses, and do a step of gradient descent.

We have developed the `autogram` engine to compute $\mathcal{G}\mathbf{f}(\mathbf{x})$ efficiently (work in progress).

This saves a lot of memory, so it's often much faster.

Gramian-based JD

$\mathcal{A}_{\text{UPGrad}}(\mathcal{J}\mathbf{f}(\mathbf{x}))$ is a linear combination of the rows of $\mathcal{J}\mathbf{f}(\mathbf{x})$.

The weights only depend on the (small) Gramian
 $\mathcal{G}\mathbf{f}(\mathbf{x}) = \mathcal{J}\mathbf{f}(\mathbf{x}) \cdot \mathcal{J}\mathbf{f}(\mathbf{x})^\top$.

We can compute $\mathcal{G}\mathbf{f}(\mathbf{x})$ directly, extract the weights, combine the losses, and do a step of gradient descent.

We have developed the `autogram` engine to compute $\mathcal{G}\mathbf{f}(\mathbf{x})$ efficiently (work in progress).

This saves a lot of memory, so it's often much faster.

Applicable to most other aggregators from the literature.